

ML_6_XGBoost

March 19, 2019

```
In [1]: """In this assignment students need to predict whether a person makes over 50K per year or not from classic adult dataset using XGBoost. The description of the dataset is as follows:"""
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
train_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data/train.csv')
test_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data/test.csv')
col_labels = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'native_country', 'wage_class']
train_set.columns = col_labels
test_set.columns = col_labels
train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
age                32561 non-null int64
workclass          32561 non-null object
fnlwgt             32561 non-null int64
education          32561 non-null object
education_num      32561 non-null int64
marital_status     32561 non-null object
occupation         32561 non-null object
relationship       32561 non-null object
race               32561 non-null object
sex                32561 non-null object
capital_gain       32561 non-null int64
capital_loss       32561 non-null int64
hours_per_week     32561 non-null int64
native_country     32561 non-null object
wage_class         32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
In [3]: # data cleaning
```

```
df = pd.concat([train_set, test_set], axis=0)
df['wage_class'] = df['wage_class'].apply(lambda x : 1 if x=='>50K' else 0) # converti
df.head()
```

```
Out[3]:
```

	age	workclass	fnlwgt	education	education_num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital_status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital_gain	capital_loss	hours_per_week	native_country	wage_class
0	2174	0	40	United-States	0
1	0	0	13	United-States	0
2	0	0	40	United-States	0
3	0	0	40	United-States	0
4	0	0	40	Cuba	0

```
In [5]: # Remove unknown values and spaces
```

```
df.replace(' ?', np.nan, inplace=True)
for col in df.columns:
    if type(df[col][0]) == str:
        print("Working on " + col)
        df[col] = df[col].apply(lambda val: val.replace(" ", ""))
df.head()
```

```
Out[5]:
```

	age	workclass	fnlwgt	education	education_num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital_status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	

	Married-civ-spouse	Prof-specialty	Wife	Black	Female
	capital_gain	capital_loss	hours_per_week	native_country	wage_class
0	2174	0	40	United-States	0
1	0	0	13	United-States	0
2	0	0	40	United-States	0
3	0	0	40	United-States	0
4	0	0	40	Cuba	0

In [7]: # Convert categorical variables into int.

```
df = pd.concat([df, pd.get_dummies(df['workclass'], prefix='workclass', prefix_sep=':')]
df.drop('workclass', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['marital_status'], prefix='marital_status', prefix_sep=':')]
df.drop('marital_status', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['occupation'], prefix='occupation', prefix_sep=':')]
df.drop('occupation', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['relationship'], prefix='relationship', prefix_sep=':')]
df.drop('relationship', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['race'], prefix='race', prefix_sep=':')]
df.drop('race', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['sex'], prefix='sex', prefix_sep=':')]
df.drop('sex', axis=1, inplace=True)

df = pd.concat([df, pd.get_dummies(df['native_country'], prefix='native_country', prefix_sep=':')]
df.drop('native_country', axis=1, inplace=True)

df.drop('education', axis=1, inplace=True)

df.head()
```

```
Out[7]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week	\
0	39	77516	13	2174	0	40	
1	50	83311	13	0	0	13	
2	38	215646	9	0	0	40	
3	53	234721	7	0	0	40	
4	28	338409	13	0	0	40	

	wage_class	workclass: Federal-gov	workclass: Local-gov	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	

```

4          0          0          0

workclass: Never-worked          ...          \
0          0          ...
1          0          ...
2          0          ...
3          0          ...
4          0          ...

native_country: Portugal native_country: Puerto-Rico \
0          0          0
1          0          0
2          0          0
3          0          0
4          0          0

native_country: Scotland native_country: South native_country: Taiwan \
0          0          0          0
1          0          0          0
2          0          0          0
3          0          0          0
4          0          0          0

native_country: Thailand native_country: Trinidad&Tobago \
0          0          0
1          0          0
2          0          0
3          0          0
4          0          0

native_country: United-States native_country: Vietnam \
0          1          0
1          1          0
2          1          0
3          1          0
4          0          0

native_country: Yugoslavia
0          0
1          0
2          0
3          0
4          0

```

[5 rows x 98 columns]

In [12]: # *Splitting into dependent and independent variables.*

```

from sklearn import preprocessing

X = np.array(df.drop(['wage_class'], 1))
y = np.array(df['wage_class'])
X = preprocessing.scale(X)      #Standard Scaling

#Splitting into train and test part

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Applying XGBoost Classifier model

import xgboost as xgb
from sklearn.metrics import accuracy_score

model = xgb.XGBClassifier(learning_rate=0.1,
                           n_estimators=500,
                           max_depth=5,
                           min_child_weight=4
                           )
model.fit(X_train, y_train)
predictions = model.predict(X_test)
XGBA = accuracy_score(y_test, predictions)
print("The Accuracy is {}".format(XGBA))

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning:
 warnings.warn(msg, DataConversionWarning)

The Accuracy is 1.0